

R package dplyr, Data Manipulation

Jason Wang

Tuesday, April 28, 2015

This is an introduction to the R package dplyr. If you type the “dplyr” in google or youtube, you will get tremendous references, tutorials and examples. It is written in C, which means that it will translate your R code into C and the manipulating functions are really fast. Its function and structure are like those we use in database language. Mainly, it contains six verb functions, filter, select, arrange, mutate, summarise, group_by and also some more advanced function.

```
library(dplyr)
#Package which I use its data to demonstrate.
library(hflights)
```

```
data(hflights)
dim(hflights)
```

```
## [1] 227496      21
```

```
names(hflights)
```

```
## [1] "Year"           "Month"           "DayofMonth"
## [4] "DayOfWeek"      "DepTime"         "ArrTime"
## [7] "UniqueCarrier"  "FlightNum"       "TailNum"
## [10] "ActualElapsedTime" "AirTime"        "ArrDelay"
## [13] "DepDelay"       "Origin"          "Dest"
## [16] "Distance"       "TaxiIn"          "TaxiOut"
## [19] "Cancelled"      "CancellationCode" "Diverted"
```

```
head(hflights)
```

If you run the head code, the output is not really good for people to understand. I am not going to show the output here.

1. Data Type: tbl

First of all, we have to transform our data into the table, local data frame. (Note that originally it is a data frame.)

```
#Transform data.frame to table
hflights <- tbl_df(hflights)
```

Normally, if you print out the data which have 227496 rows and 21 columns, you will get a messy output. However, when the data is table, then we can print it without any worry. It will automatically adapt the data output to us.

```
#Nicer output compared with the originally one
hflights
```

```
## Source: local data frame [227,496 x 21]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     1           1         6   1400   1500           AA         428
## 2  2011     1           2         7   1401   1501           AA         428
## 3  2011     1           3         1   1352   1502           AA         428
## 4  2011     1           4         2   1403   1513           AA         428
## 5  2011     1           5         3   1405   1507           AA         428
## 6  2011     1           6         4   1359   1503           AA         428
## 7  2011     1           7         5   1359   1509           AA         428
## 8  2011     1           8         6   1355   1454           AA         428
## 9  2011     1           9         7   1443   1554           AA         428
## 10 2011     1          10         1   1443   1553           AA         428
## .. ... ..
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
## (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
## Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
## CancellationCode (chr), Diverted (int)
```

In the beginning, it shows you the dimension of the data. In the middle, it gives you an appropriate output which fit the windows. In the end, it tells you what variables do not show.

2. Verb 1 filter (Row)

Usually, when we first get the data, we will explore it in a different way. We may think that maybe certain group in the data will have the same patterns. We want to find out those subjects (row) which have the same features in certain variable. Then, we can utilize the `filter` function to help us.

For instance, you may want to get those flights in January. In `dplyr`, it is really easy.

```
#In dplyr
filter(hflights, Month == 1)

## Source: local data frame [18,910 x 21]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     1           1         6    1400    1500           AA         428
## 2  2011     1           2         7    1401    1501           AA         428
## 3  2011     1           3         1    1352    1502           AA         428
## 4  2011     1           4         2    1403    1513           AA         428
## 5  2011     1           5         3    1405    1507           AA         428
## 6  2011     1           6         4    1359    1503           AA         428
## 7  2011     1           7         5    1359    1509           AA         428
## 8  2011     1           8         6    1355    1454           AA         428
## 9  2011     1           9         7    1443    1554           AA         428
##10  2011     1          10         1    1443    1553           AA         428
## .. ... ..
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
## (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
## Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
## CancellationCode (chr), Diverted (int)
```

As you can see, it will show you a nicely format. You can know how many flights are in January. In normal R codes, you may have to write some codes like `hflights[hflights$Month == 1,]`.

You can save it into another table such as `Jan`. It will keep the same data type. Also, you can transform it back to `data.frame` without any warning message.

```
Jan <- filter(hflights, Month == 1)
class(Jan)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
#You can also get the data as the way for data frame
Jan[1, 1]
```

```
## Source: local data frame [1 x 1]
##
##   Year
## 1 2011
```

```
Jan <- data.frame(Jan)
class(Jan)
```

```
## [1] "data.frame"
```

If you have multiple condition, you can use AND(&) as well as OR(|) in the function.

```
filter(hflights, Month == 1 & DayofMonth == 1)
```

```
## Source: local data frame [552 x 21]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     1           1           6    1400    1500           AA         428
## 2  2011     1           1           6     728     840           AA         460
## 3  2011     1           1           6    1631    1736           AA        1121
## 4  2011     1           1           6    1756    2112           AA        1294
## 5  2011     1           1           6    1012    1347           AA        1700
## 6  2011     1           1           6    1211    1325           AA        1820
## 7  2011     1           1           6     557     906           AA        1994
## 8  2011     1           1           6    1824    2106           AS         731
## 9  2011     1           1           6     654    1124           B6         620
## 10 2011     1           1           6    1639    2110           B6         622
## .. ... ..
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
##   (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
##   Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
##   CancellationCode (chr), Diverted (int)
```

```
#You can also replace AND by ,
filter(hflights, Month == 1, DayofMonth == 1)
```

```
## Source: local data frame [552 x 21]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     1           1           6    1400    1500           AA         428
## 2  2011     1           1           6     728     840           AA         460
## 3  2011     1           1           6    1631    1736           AA        1121
```

```
## 4 2011      1      1      6    1756    2112      AA    1294
## 5 2011      1      1      6    1012    1347      AA    1700
## 6 2011      1      1      6    1211    1325      AA    1820
## 7 2011      1      1      6     557     906      AA    1994
## 8 2011      1      1      6    1824    2106      AS     731
## 9 2011      1      1      6     654    1124      B6     620
## 10 2011     1      1      6    1639    2110      B6     622
## .. ... .. ... .. ... .. ... ..
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
## (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
## Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
## CancellationCode (chr), Diverted (int)
```

You can also filter by character.

```
filter(hflights, UniqueCarrier == "AA")
```

```
## Source: local data frame [3,244 x 21]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     1          1          6    1400    1500          AA         428
## 2  2011     1          2          7    1401    1501          AA         428
## 3  2011     1          3          1    1352    1502          AA         428
## 4  2011     1          4          2    1403    1513          AA         428
## 5  2011     1          5          3    1405    1507          AA         428
## 6  2011     1          6          4    1359    1503          AA         428
## 7  2011     1          7          5    1359    1509          AA         428
## 8  2011     1          8          6    1355    1454          AA         428
## 9  2011     1          9          7    1443    1554          AA         428
## 10 2011     1         10          1    1443    1553          AA         428
## .. ... .. ... .. ... .. ... ..
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
## (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
## Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
## CancellationCode (chr), Diverted (int)
```

3. Verb 2 select (Column)

You might also want to examine some column in the data. In R, you may use the code like `hflights[, c("Month", "DayofMonth", "FlightNum")]`.

```
#In dplyr  
select(hflights, Month, DayofMonth, FlightNum)
```

```
## Source: local data frame [227,496 x 3]  
##  
##   Month DayofMonth FlightNum  
## 1      1          1        428  
## 2      1          2        428  
## 3      1          3        428  
## 4      1          4        428  
## 5      1          5        428  
## 6      1          6        428  
## 7      1          7        428  
## 8      1          8        428  
## 9      1          9        428  
## 10     1         10        428  
## ..      ...          ...          ...
```

```
#You can also use contain function to select the column which have the same keyword  
select(hflights, contains("delay"))
```

```
## Source: local data frame [227,496 x 2]  
##  
##   ArrDelay DepDelay  
## 1      -10         0  
## 2       -9         1  
## 3       -8        -8  
## 4         3         3  
## 5       -3         5  
## 6       -7        -1  
## 7       -1        -1  
## 8      -16        -5  
## 9        44        43  
## 10       43        43  
## ..      ...          ...
```

```
#Or use starts_with
select(hflights, starts_with("M"))
```

```
## Source: local data frame [227,496 x 1]
##
##   Month
## 1      1
## 2      1
## 3      1
## 4      1
## 5      1
## 6      1
## 7      1
## 8      1
## 9      1
## 10     1
## ..    ...
```

```
#Or use ends_with
select(hflights, ends_with("th"))
```

```
## Source: local data frame [227,496 x 2]
##
##   Month DayofMonth
## 1      1          1
## 2      1          2
## 3      1          3
## 4      1          4
## 5      1          5
## 6      1          6
## 7      1          7
## 8      1          8
## 9      1          9
## 10     1         10
## ..    ...        ...
```

Digression about a nicer way to write codes in dplyr

Sometimes we will combine the `filter` and `select` function together. For example, you may want to find out the FlightNum and Origin of those flights which their UniqueCarrier are WN. You will write a code like the following:

```
filter(select(hflights, FlightNum, Origin, UniqueCarrier), UniqueCarrier == "WN")
```

```
## Source: local data frame [45,343 x 3]
```

```
##
```

```
##   FlightNum Origin UniqueCarrier
## 1      1266    HOU              WN
## 2      1689    HOU              WN
## 3      1024    HOU              WN
## 4      2430    HOU              WN
## 5      3013    HOU              WN
## 6      1038    HOU              WN
## 7      2345    HOU              WN
## 8      1454    HOU              WN
## 9      2360    HOU              WN
## 10     1593    HOU              WN
## ..      ...      ...          ...
```

When you see the above codes, it may take some times to really get what the codes is doing. However, in dplyr, it gives user a more generally way to organize their code called “**chaining**” which can give a more readable code.

```
#Another much more easier way to write code
```

```
hflights %>%
  select(FlightNum, Origin, UniqueCarrier) %>%
  filter(UniqueCarrier == "WN")
```

```
## Source: local data frame [45,343 x 3]
```

```
##
```

```
##   FlightNum Origin UniqueCarrier
## 1      1266    HOU              WN
## 2      1689    HOU              WN
## 3      1024    HOU              WN
## 4      2430    HOU              WN
## 5      3013    HOU              WN
## 6      1038    HOU              WN
## 7      2345    HOU              WN
## 8      1454    HOU              WN
```



```
## 9      2360    HOU      WN
## 10     1593    HOU      WN
## ..      ...      ...      ...
```

You can regard the %>% syntax as “then”. It will be like that get the hflights data, then select FlightNum, Origin, UniqueCarrier, then filter the outcome to find those UniqueCarrier is “WN”. It is much more interpretable way to present the code to others. Actually, the syntax exists in the R code and we can utilize it in other code too.

```
x <- rnorm(2)
y <- rnorm(2)
#Distance
sqrt(sum((x - y)^2))
```

```
## [1] 3.098904
```

```
#In %>% way
(x - y)^2 %>% sum %>% sqrt
```

```
## [1] 3.098904
```

4. Verb 3 arrange

Usually, we want to sort data and see whether overall data shows pattern after sorting. We can use `hflights[order(hflights$DepTime), c("Month", "DepTime", "ArrTime")]`

```
#In dplyr
arrange(select(hflights, Month, DepTime, ArrTime), DepTime)
```

```
## Source: local data frame [227,496 x 3]
```

```
##
```

```
##   Month DepTime ArrTime
## 1     1       1      621
## 2     3       1      557
## 3     4       1      510
## 4     6       1      515
## 5    11       1       55
## 6    12       1      642
## 7    12       1      633
## 8     7       2       53
## 9    12       2      611
## 10    7       3      521
## ..      ...      ...      ...
```

#Chaining

```
hflights %>%  
  select(Month, DepTime, ArrTime) %>%  
  arrange(DepTime)
```

```
## Source: local data frame [227,496 x 3]
```

```
##
```

```
##   Month DepTime ArrTime
```

```
## 1     1         1     621
```

```
## 2     3         1     557
```

```
## 3     4         1     510
```

```
## 4     6         1     515
```

```
## 5    11         1      55
```

```
## 6    12         1     642
```

```
## 7    12         1     633
```

```
## 8     7         2      53
```

```
## 9    12         2     611
```

```
## 10    7         3     521
```

```
## .. ...      ...      ...
```

#In decreasing way

```
hflights %>%  
  select(Month, DepTime, ArrTime) %>%  
  arrange(desc(DepTime))
```

```
## Source: local data frame [227,496 x 3]
```

```
##
```

```
##   Month DepTime ArrTime
```

```
## 1     5     2400     144
```

```
## 2     4     2359     455
```

```
## 3     5     2359     130
```

```
## 4     5     2359      56
```

```
## 5     6     2359     113
```

```
## 6     6     2359      40
```

```
## 7     6     2359     111
```

```
## 8     7     2359     108
```

```
## 9     7     2359     105
```

```
## 10    9     2359     106
```

```
## .. ...      ...      ...
```

5. Verb 4 mutate

Sometimes, you may want to add or create a new variable into data. For example, if we want to add speed, we use `hflights$Speed <- hflights$Distance/hflights$AirTime`.

```
#In dplyr
hflights %>%
  select(Distance, AirTime) %>%
  mutate(Speed=Distance/AirTime)
```

```
## Source: local data frame [227,496 x 3]
##
##   Distance AirTime   Speed
## 1      224      40 5.600000
## 2      224      45 4.977778
## 3      224      48 4.666667
## 4      224      39 5.743590
## 5      224      44 5.090909
## 6      224      45 4.977778
## 7      224      43 5.209302
## 8      224      40 5.600000
## 9      224      41 5.463415
## 10     224      45 4.977778
## ..      ...      ...      ...
```

```
#To store
hflights <- mutate(hflights, Speed=Distance/AirTime)
select(hflights, Distance, AirTime, Speed)
```

```
## Source: local data frame [227,496 x 3]
##
##   Distance AirTime   Speed
## 1      224      40 5.600000
## 2      224      45 4.977778
## 3      224      48 4.666667
## 4      224      39 5.743590
## 5      224      44 5.090909
## 6      224      45 4.977778
## 7      224      43 5.209302
## 8      224      40 5.600000
## 9      224      41 5.463415
## 10     224      45 4.977778
## ..      ...      ...      ...
```

6. Verb 5 summarise + Verb 6 group_by

When it comes to descriptive statistics, we will want to statistics based on several group if we have multiple group. For example, if we want to know the average arrival delay time for different destination, then we run `head(tapply(hflights$ArrDelay, hflights$Dest, mean, na.rm=T))` or `head(aggregate(ArrDelay ~ Dest, hflights, mean))`.

#In dplyr

```
summarise(group_by(hflights, Dest), mean(ArrDelay, na.rm=T))
```

```
## Source: local data frame [116 x 2]
##
##   Dest mean(ArrDelay, na.rm = T)
## 1  ABQ          7.226259
## 2  AEX          5.839437
## 3  AGS          4.000000
## 4  AMA          6.840095
## 5  ANC         26.080645
## 6  ASE          6.794643
## 7  ATL          8.233251
## 8  AUS          7.448718
## 9  AVL          9.973988
## 10 BFL         -13.198807
## .. ...
```

#Chaining

```
hflights %>%
  group_by(Dest) %>%
  summarise(delay_Time=mean(ArrDelay, na.rm=T))
```

```
## Source: local data frame [116 x 2]
##
##   Dest delay_Time
## 1  ABQ    7.226259
## 2  AEX    5.839437
## 3  AGS    4.000000
## 4  AMA    6.840095
## 5  ANC   26.080645
## 6  ASE    6.794643
## 7  ATL    8.233251
## 8  AUS    7.448718
## 9  AVL    9.973988
## 10 BFL  -13.198807
## .. ...
```

Sometimes, we may want to apply the function to several columns or apply different functions to column.

```
#Apply function to multiple columns
```

```
hflights %>%  
  group_by(UniqueCarrier) %>%  
  summarise_each(funs(mean(., na.rm=T)), Cancelled, Diverted)
```

```
## Source: local data frame [15 x 3]  
##  
##   UniqueCarrier  Cancelled  Diverted  
## 1           AA 0.018495684 0.001849568  
## 2           AS 0.000000000 0.002739726  
## 3           B6 0.025899281 0.005755396  
## 4           C0 0.006782614 0.002627370  
## 5           DL 0.015903067 0.003029156  
## 6           EV 0.034482759 0.003176044  
## 7           F9 0.007159905 0.000000000  
## 8           FL 0.009817672 0.003272557  
## 9           MQ 0.029044750 0.001936317  
## 10          00 0.013946828 0.003486707  
## 11          UA 0.016409266 0.002413127  
## 12          US 0.011268986 0.001469868  
## 13          WN 0.015504047 0.002293629  
## 14          XE 0.015495599 0.003449550  
## 15          YV 0.012658228 0.000000000
```

```
#Apply multiple functions to one column
```

```
hflights %>%  
  group_by(Dest) %>%  
  summarise_each(funs(mean(., na.rm=T), min(., na.rm=T), max(., na.rm=T)), ArrDelay)
```

```
## Source: local data frame [116 x 4]  
##  
##   Dest      mean min max  
## 1  ABQ    7.226259 -26 290  
## 2  AEX    5.839437 -34 257  
## 3  AGS    4.000000  4   4  
## 4  AMA    6.840095 -28 301  
## 5  ANC   26.080645 -21 281  
## 6  ASE    6.794643 -31 252  
## 7  ATL    8.233251 -41 701  
## 8  AUS    7.448718 -24 244  
## 9  AVL    9.973988 -23 331  
## 10 BFL  -13.198807 -56 206  
## .. ...      ... ..
```

```
#Apply multiple functions to several columns
```

```
hflights %>%
```

```
  group_by(UniqueCarrier) %>%
```

```
  summarise_each(funs(min(., na.rm=T), max(., na.rm=T)), contains("Delay"))
```

```
## Source: local data frame [15 x 5]
```

```
##
```

##	UniqueCarrier	ArrDelay_min	DepDelay_min	ArrDelay_max	DepDelay_max
## 1	AA	-39	-15	978	970
## 2	AS	-43	-15	183	172
## 3	B6	-44	-14	335	310
## 4	CO	-55	-18	957	981
## 5	DL	-32	-17	701	730
## 6	EV	-40	-18	469	479
## 7	F9	-24	-15	277	275
## 8	FL	-30	-14	500	507
## 9	MQ	-38	-23	918	931
## 10	OO	-57	-33	380	360
## 11	UA	-47	-11	861	869
## 12	US	-42	-17	433	425
## 13	WN	-44	-10	499	548
## 14	XE	-70	-19	634	628
## 15	YV	-32	-11	72	54

7. n() and n_distinct

Here we introduce a useful functions that can facilitate our analysis. The first one is `n()`. It can help us count the number in the group(category) and perhaps we want to sort them by the its numbers. For example, we may want to know what's the rush month or day for airline. We can count the number by these two category and sort them to find out the results.

```
#In dplyr
```

```
hflights %>%
```

```
  group_by(Month, DayofMonth) %>%
```

```
  summarise(flight_count=n()) %>%
```

```
  ungroup() %>%
```

```
  arrange(desc(flight_count))
```

```
## Source: local data frame [365 x 3]
```

```
##
```

##	Month	DayofMonth	flight_count
## 1	8	4	706

```
## 2      8      11      706
## 3      8      12      706
## 4      8       5      705
## 5      8       3      704
## 6      8      10      704
## 7      1       3      702
## 8      7       7      702
## 9      7      14      702
## 10     7      28      701
## ..     ...     ...     ...
```

```
#Another easier way to write it
hflights %>%
  group_by(Month, DayofMonth) %>%
  tally(sort=T)
```

```
## Source: local data frame [365 x 3]
## Groups: Month
##
##   Month DayofMonth   n
## 1     1           3 702
## 2     1           2 678
## 3     1          20 663
## 4     1          27 663
## 5     1          13 662
## 6     1           7 661
## 7     1          14 661
## 8     1          21 661
## 9     1          28 661
## 10    1           6 660
## ..     ...     ... ..
```

`n_distinct` function provide the number of unique subject in certain column you assign.

```
hflights %>%
  group_by(Dest) %>%
  summarise(flight_count=n(), plane_count=n_distinct(TailNum))
```

```
## Source: local data frame [116 x 3]
##
##   Dest flight_count plane_count
## 1  ABQ          2812          716
## 2  AEX           724          215
```

```
## 3   AGS           1           1
## 4   AMA         1297         158
## 5   ANC          125          38
## 6   ASE          125          60
## 7   ATL         7886         983
## 8   AUS         5022        1015
## 9   AVL          350         142
## 10  BFL          504          70
## .. ...           ...           ...
```

Sometimes, `group_by` can be very useful even without summarise.

```
hflights %>%
  group_by(Dest) %>%
  select(Cancelled) %>%
  table() %>%
  head()
```

```
##      Cancelled
## Dest      0   1
## ABQ 2787 25
## AEX  712 12
## AGS    1  0
## AMA 1265 32
## ANC  125  0
## ASE  120  5
```

8. Windows function

- Aggregate function: Input n values, output 1 value

*Windows function: Input n values, output n values

*#For each destination, calculate which two days of the year they had their longest dep
#We use desc(DepDelay) because I want to find the largest value of DepDelay.*

```
hflights %>%
  group_by(Dest) %>%
  select(Month, DayofMonth, DepDelay) %>%
  filter(min_rank(desc(DepDelay)) <= 2) %>%
  arrange(Dest, desc(DepDelay)) %>%
  print(n=15)
```



```
## Source: local data frame [233 x 4]
## Groups: Dest
##
##   Dest Month DayofMonth DepDelay
## 1  ABQ     3         31      300
## 2  ABQ    10         24      275
## 3  AEX    12         31      266
## 4  AEX     2          5      173
## 5  AGS     4          3       10
## 6  AMA    10          9      304
## 7  AMA    10          9      246
## 8  ANC     8         11      292
## 9  ANC     6         25      215
## 10 ASE    12         29      269
## 11 ASE     2          7      208
## 12 ATL    10         25      730
## 13 ATL     2         19      507
## 14 AUS     6         22      240
## 15 AUS    10          9      239
## .. ... ..
```

- top_n

```
#top_n
hflights %>%
  group_by(Dest) %>%
  select(Month, DayofMonth, DepDelay) %>%
  top_n(2) %>%
  arrange(Dest, desc(DepDelay))
```

Selecting by DepDelay

```
## Source: local data frame [233 x 4]
## Groups: Dest
##
##   Dest Month DayofMonth DepDelay
## 1  ABQ     3         31      300
## 2  ABQ    10         24      275
## 3  AEX    12         31      266
## 4  AEX     2          5      173
## 5  AGS     4          3       10
## 6  AMA    10          9      304
## 7  AMA    10          9      246
## 8  ANC     8         11      292
```

```
## 9   ANC      6      25      215
## 10  ASE     12      29      269
## ..   ...     ...     ...     ...
```

```
#Compute the change from month to month
hflights %>%
  group_by(Month) %>%
  summarise(flight_count=n()) %>%
  mutate(change=flight_count-lag(flight_count))
```

```
## Source: local data frame [12 x 3]
##
##   Month flight_count change
## 1     1      18910      NA
## 2     2      17128    -1782
## 3     3      19470     2342
## 4     4      18593     -877
## 5     5      19172      579
## 6     6      19600      428
## 7     7      20548      948
## 8     8      20176     -372
## 9     9      18065    -2111
## 10    10      18696      631
## 11    11      18021     -675
## 12    12      19117     1096
```

```
#By tally function
hflights %>%
  group_by(Month) %>%
  tally() %>%
  mutate(change=n - lag(n))
```

```
## Source: local data frame [12 x 3]
##
##   Month      n change
## 1     1 18910      NA
## 2     2 17128    -1782
## 3     3 19470     2342
## 4     4 18593     -877
## 5     5 19172      579
## 6     6 19600      428
## 7     7 20548      948
## 8     8 20176     -372
```

```
## 9      9 18065 -2111
## 10     10 18696   631
## 11     11 18021  -675
## 12     12 19117  1096
```

9. Others

We can also do sampling easily.

```
#In dplyr
hflights %>% sample_n(10)
```

```
## Source: local data frame [10 x 22]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011     6           5          7   1909    2004           XE       3027
## 2  2011     8           5          5   1122    1240           XE       2277
## 3  2011    11          14          1   1445    1819           CO       1623
## 4  2011     7          21          4    758    1104           XE       2470
## 5  2011    12          17          6    722     920           XE       4676
## 6  2011     4          24          7   1910    2304           XE       2171
## 7  2011     8          18          4   1220    1316           CO       1555
## 8  2011     3          10          4   1551    1734           OO       1173
## 9  2011    12           9          5   1302    1643           CO       1653
## 10 2011     7           4          1   1152    1255           CO       1629
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
##   (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
##   Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
##   CancellationCode (chr), Diverted (int), Speed (dbl)
```

```
#Or by fraction
hflights %>% sample_frac(0.25, replace=T)
```

```
## Source: local data frame [56,874 x 22]
##
##   Year Month DayofMonth DayOfWeek DepTime ArrTime UniqueCarrier FlightNum
## 1  2011    11          18          5   1323    1520           OO       5249
## 2  2011    11          30          3    807    1126           CO       1160
## 3  2011     1          26          3   1141    1412           CO        546
## 4  2011     2          27          7   1114    1353           WN         11
## 5  2011     1          11          2    729    1026           XE      2586
## 6  2011     1           6          4   1907    2029           XE      2280
```

```
## 7 2011      3      28      1    1036    1610      CO      212
## 8 2011      8      25      4    1436    1525      XE      2451
## 9 2011      8      15      1    2056    2119      OO      1108
## 10 2011     4      22      5    2128    2235      WN      776
## .. ... .. ... .. ... .. ...
## Variables not shown: TailNum (chr), ActualElapsedTime (int), AirTime
## (int), ArrDelay (int), DepDelay (int), Origin (chr), Dest (chr),
## Distance (int), TaxiIn (int), TaxiOut (int), Cancelled (int),
## CancellationCode (chr), Diverted (int), Speed (dbl)
```

```
#Like the str function in base R
glimpse(hflights)
```

```
## Observations: 227496
## Variables:
## $ Year      (int) 2011, 2011, 2011, 2011, 2011, 2011, 2011, 20...
## $ Month     (int) 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,...
## $ DayofMonth (int) 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 1...
## $ DayOfWeek (int) 6, 7, 1, 2, 3, 4, 5, 6, 7, 1, 2, 3, 4, 5, 6,...
## $ DepTime   (int) 1400, 1401, 1352, 1403, 1405, 1359, 1359, 13...
## $ ArrTime   (int) 1500, 1501, 1502, 1513, 1507, 1503, 1509, 14...
## $ UniqueCarrier (chr) "AA", "AA", "AA", "AA", "AA", "AA", "AA", "A...
## $ FlightNum (int) 428, 428, 428, 428, 428, 428, 428, 428, 428,...
## $ TailNum   (chr) "N576AA", "N557AA", "N541AA", "N403AA", "N49...
## $ ActualElapsedTime (int) 60, 60, 70, 70, 62, 64, 70, 59, 71, 70, 70, ...
## $ AirTime   (int) 40, 45, 48, 39, 44, 45, 43, 40, 41, 45, 42, ...
## $ ArrDelay  (int) -10, -9, -8, 3, -3, -7, -1, -16, 44, 43, 29,...
## $ DepDelay  (int) 0, 1, -8, 3, 5, -1, -1, -5, 43, 43, 29, 19, ...
## $ Origin    (chr) "IAH", "IAH", "IAH", "IAH", "IAH", "IAH", "I...
## $ Dest      (chr) "DFW", "DFW", "DFW", "DFW", "DFW", "DFW", "D...
## $ Distance  (int) 224, 224, 224, 224, 224, 224, 224, 224, 224,...
## $ TaxiIn    (int) 7, 6, 5, 9, 9, 6, 12, 7, 8, 6, 8, 4, 6, 5, 6...
## $ TaxiOut   (int) 13, 9, 17, 22, 9, 13, 15, 12, 22, 19, 20, 11...
## $ Cancelled (int) 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ CancellationCode (chr) "", "", "", "", "", "", "", "", "", "", "", ...
## $ Diverted  (int) 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ Speed     (dbl) 5.600000, 4.977778, 4.666667, 5.743590, 5.09...
```

10. Connecting with database

- dplyr can connect to a database as if the data was loaded into a data frame
- [Instruction for create a database](#)

```

# connect to an SQLite database
my_db <- src_sqlite("my_db.sqlite3")

# connect to the "hflights" table in that database
flights_tbl <- tbl(my_db, "hflights")

# identical query using the database
flights_tbl %>%
  select(UniqueCarrier, DepDelay) %>%
  arrange(desc(DepDelay))

# ask dplyr for the SQL commands
flights_tbl %>%
  select(UniqueCarrier, DepDelay) %>%
  arrange(desc(DepDelay)) %>%
  explain()

```

References

- [Hands-on dplyr tutorial for faster data manipulation in R](#)
- [A more comprehensive and advanced tutorial](#)